

# PERN v4.0

## AI TRUST & PREDICTION ENGINE

Machine Learning Technical Documentation — features, data sources, data volume, methodology & technology

Served artifact: persistence · trained 2026-08-12T02:15:49 UTC · generated 2026-08-14 13:59 UTC

1.23 °C

Temporal RMSE

90.4%

Temporal coverage

3.40 °C

Interval width (90%)

87.9


Confidence score

## 1 · Overview

The PERN AI Trust & Prediction Engine is a machine-learning subsystem that turns environmental observations into a **calibrated prediction interval** and a 0–100 **confidence score**. It is deliberately decoupled from the website until validated; the backend reaches it through a fail-open HTTP client (4 s timeout, 60 s cache).

Core insight: at a 1-day forecasting horizon, daily-mean temperature is persistence-dominated (today ≈ tomorrow). The served model therefore uses **persistence as the point forecast** and applies **split-conformal calibration to the residuals** so interval coverage is guaranteed at the declared level. The engine owns the *uncertainty*, not a better point forecast.

## 2 · Pipeline



Air track flows identically from OpenAQ (or its simulated fallback) through the same feature/label/train/serve stages.

## 3 · Data Sources

|                                     |                                                                                                                                                                                                                               |
|-------------------------------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>NASA POWER (keyless)</b>         | power.larc.nasa.gov temporal daily API · community AG · no key. Live-fetched per grid point (start/end as YYYYMMDD), JSON response, -999 missing sentinel cleaned.                                                            |
| <b>OpenAQ (key-gated)</b>           | openaq.fetchByLocation(lat, lng, limit) adapter wired in (openaq-source.js:64). Real key + daily labels still required for a production air model; the air track is exercised today with synthetic labels in the same schema. |
| <b>PostgreSQL (feature_vectors)</b> | feature-etl.js writes aligned (location × time) rows with feature_group, source_id, snapshot, quality and provenance; the training pipeline reads them back and attaches labels.                                              |

NASA POWER parameters fetched (agriculture track): **T2M**, **T2M\_MAX**, **T2M\_MIN**, **PRECTOTCORR** (corrected precipitation), **RH2M** (relative humidity), **WS2M** (2-m wind speed). All values are daily means/aggregates.

## 4 · Feature Catalog

Agriculture track (served). One row per grid point per day; label = next day's observed T2M.

| Feature             | Type  | Source            | Role                                    |
|---------------------|-------|-------------------|-----------------------------------------|
| temperature         | float | POWER T2M         | persistence center + conformal residual |
| temperature_max     | float | POWER T2M_MAX     | calibration context                     |
| temperature_min     | float | POWER T2M_MIN     | calibration context                     |
| precipitation       | float | POWER PRECTOTCORR | calibration context                     |
| humidity            | float | POWER RH2M        | calibration context                     |
| wind_speed          | float | POWER WS2M        | calibration context                     |
| latitude, longitude | float | grid / request    | site identity                           |
| month               | int   | derived           | seasonal context                        |
| day_of_year         | int   | derived           | seasonal context                        |
| day_of_week         | int   | derived           | weekly context                          |

Air track (pipeline-ready, synthetic labels). Label = next-day PM2.5.

| Feature                           | Type                       |
|-----------------------------------|----------------------------|
| temperature, humidity, wind_speed | float · POWER meteorology  |
| pm25, pm10                        | float · particulate        |
| no2, so2, o3                      | float · gaseous pollutants |
| month, day_of_year, day_of_week   | int · derived              |

## 5 · Data Volume

|                                      |                                                                                                                                                                                                                                                                          |
|--------------------------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>Real agriculture dataset</b>      | 1,888 labeled rows · 16 grid points (4×4 lat/lon grid) · ~120 days (2026-04-12 → 2026-08-07) · one (site, day) row with next-day T2M label                                                                                                                               |
| <b>Served-artifact training rows</b> | 1,888 · feature columns: 11                                                                                                                                                                                                                                              |
| <b>Air dataset</b>                   | 1,904 rows · 16 sites · 120 days — physically-plausible PM2.5 (traffic weekday cycle, wind dilution, humidity, site emission); fetch_openaq_labels.py replaces them with real OpenAQ daily labels once OPENAQ_API_KEY is set (fallback keeps the pipeline offline-green) |
| <b>Synthetic fallback</b>            | app/ml/synthetic.py reproduces the exact schema for offline dev / smoke tests when no dataset exists                                                                                                                                                                     |
| <b>Splits</b>                        | Temporal backtest ~1,416 test rows across folds · spatial leave-locations-out · final holdout 240 rows (14 most-recent days)                                                                                                                                             |

## 6 · ML Methodology & Model Spec

|                              |                                                                                                                                                                                                                                                                  |
|------------------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>Served model</b>          | PersistenceResidual (app/ml/models.py): center = today's temperature ( $X[:, \text{temp\_idx}]$ ); predict_interval returns zero-width [center, center]; the conformal layer adds $\pm \text{qhat}$ . Zero fitted parameters — the honest optimum for this task. |
| <b>Tree alternatives</b>     | LightGBMQuantile (7 quantile regressors: 0.05–0.95) and SklearnQuantile are implemented for other targets; they lose to persistence on real temperature at 1-day horizon.                                                                                        |
| <b>Conformal calibration</b> | split-conformal CQR: residuals $r =  y - \text{center} $ on the calibration slice; $\text{qhat} = Q_{\lfloor (n+1)(1-\alpha)/n \rfloor}(r)$ , $\alpha = 0.10$ ; interval = [center – qhat, center + qhat].                                                       |
| <b>CQR (tree path)</b>       | scores = max(lower – y, y – upper) from quantile models at $\alpha/2$ and $1-\alpha/2$ ; same qhat formula, applied as [lower – qhat, upper + qhat].                                                                                                             |


|                         |                                                                                                                                                                                                                                                                                                                                                                             |
|-------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>Served model</b>     | PersistenceResidual (app/ml/models.py): center = today's temperature ( $X[:, \text{temp\_idx}]$ ); predict_interval returns zero-width [center, center]; the conformal layer adds $\pm \hat{q}$ . Zero fitted parameters — the honest optimum for this task.                                                                                                                |
| <b>Confidence score</b> | coverage_match = $1 -  \text{coverage} - (1 - \alpha) $ ; width_score = $1 / (1 + \text{width} / (2 \cdot \text{target\_std}))$ — a monotonic, non-saturating mapping so any interval widening (adaptive factors for new sites / off-distribution inputs) always lowers the score; score = $100 \cdot (0.6 \cdot \text{coverage\_match} + 0.4 \cdot \text{width\_score})$ . |
| <b>Anti-leakage</b>     | temporal_block_split (train strictly precedes test), leave_locations_out_split (whole sites held out), calibration carved from the most recent 30% of each train block, NaN-imputation medians computed from training rows only, and a final 14-day holdout never touched by model selection.                                                                               |

## 6.1 · Hyperparameters


| Parameter                   | Value                                                                                   | Scope                                                       |
|-----------------------------|-----------------------------------------------------------------------------------------|-------------------------------------------------------------|
| alpha (conformal)           | 0.10                                                                                    | target 90% marginal coverage                                |
| calib_frac                  | 0.30                                                                                    | most-recent fraction of each train block used for $\hat{q}$ |
| $\hat{q}$ level             | $\lfloor (n+1)(1-\alpha)/n \rfloor$                                                     | finite-sample correction in conformal.py                    |
| alphas (trees)              | (0.05, 0.10, 0.25, 0.50, 0.75, 0.90, 0.95)                                              | quantile grid / CRPS                                        |
| n_estimators (LightGBM)     | 200 (100 in --fast)                                                                     | per quantile level                                          |
| learning_rate (LightGBM)    | 0.05                                                                                    |                                                             |
| num_leaves (LightGBM)       | 31 (24 in experiments)                                                                  |                                                             |
| min_data_in_leaf (LightGBM) | 20                                                                                      |                                                             |
| objective                   | quantile                                                                                | LightGBM per-alpha regressors                               |
| max_iter (SklearnQuantile)  | 150                                                                                     | fallback path                                               |
| PSI thresholds              | moderate > 0.10 · severe > 0.25                                                         | drift                                                       |
| CUSUM                       | $k = 0.5\sigma$ · $h = 5\sigma$                                                         | drift on residuals                                          |
| serve-time drift            | $ z  > 4$ flags a feature                                                               | per request vs reference_stats                              |
| promotion gate              | coverage within $\pm 0.02$ of nominal · $\text{RMSE} \leq 1.02 \times \text{incumbent}$ | retrain.py                                                  |

## 7 · Evaluation Results (real data, served model)


| Split                | RMSE °C | MAE   | Coverage | Width °C | Conf. |
|----------------------|---------|-------|----------|----------|-------|
| Temporal (forecast)  | 1.234   | 0.781 | 90.4%    | 3.399    | 87.9  |
| Spatial (new sites)  | 1.433   | 0.954 | 77.2%    | 2.522    | 82.8  |
| Final holdout (14 d) | 0.816   | 0.607 | 87.5%    | 2.262    | 89.7  |



Model-selection experiment: \*final holdout RMSE (persistence + conformal). Lag + cross-site neighbor features and delta transforms could not beat persistence; the unshrunk delta tree was the worst tree variant.



Coverage vs the 90% target. The spatial split is intentionally harder — unseen sites swing more, and the engine correctly refuses to overstate confidence there. The final holdout (240 rows, never used for selection) sits within binomial noise of nominal.



Distribution of 1-day persistence residuals over the full real dataset (1,888 rows). The conformal  $\hat{q}$  (1.26 °C) brackets ~90% of the mass by construction.

## 7.1 · Multi-horizon Forecast Skill (24h / 7d / 30d)

Forecast skill as the horizon grows. Two centers are compared per horizon on both tracks: **persistence** (center = today's observation, the served model's choice) and a **seasonal climatology** (smoothed day-of-year mean fit on the calibration slice). For each site the forecast center for day  $t+h$  is formed from data up to day  $t$ , the interval


width is calibrated on the  $|\text{actual} - \text{center}|$  residuals of the **first 60% of each site's timeline** (grouped per month, split-conformal,  $\alpha = 0.10$ ), and the **last 40% is scored against real observations** — no test leakage. Best center per horizon is highlighted (gold) and marked \* in the chart. Tolerance accuracy is the competitor-style metric providers actually publish (share of forecasts within  $\pm x$ ).

#### Agriculture — NASA POWER daily T2M (°C), 16 grid sites

| Horizon | Center      | RMSE °C | MAE  | Coverage | Width | Conf. |
|---------|-------------|---------|------|----------|-------|-------|
| 1d      | persistence | 0.76    | 0.56 | 99.9%    | 6.34  | 68.5  |
| 1d      | climatology | 3.50    | 3.03 | 73.0%    | 8.22  | 61.9  |
| 7d      | persistence | 1.68    | 1.36 | 98.6%    | 11.72 | 63.9  |
| 7d      | climatology | 3.02    | 2.58 | 81.2%    | 8.31  | 66.5  |
| 30d     | persistence | 2.51    | 2.08 | 86.5%    | 9.48  | 68.4  |
| 30d     | climatology | 1.88    | 1.56 | 87.5%    | 5.90  | 73.1  |

#### Air — PM2.5 virtual-sensor track (µg/m³), 16 sites

| Horizon | Center      | RMSE µg/m³ | MAE   | Coverage | Width | Conf. |
|---------|-------------|------------|-------|----------|-------|-------|
| 1d      | persistence | 19.92      | 14.98 | 88.0%    | 62.26 | 74.1  |
| 1d      | climatology | 20.79      | 14.78 | 86.6%    | 57.45 | 74.0  |
| 7d      | persistence | 17.20      | 13.08 | 91.7%    | 60.98 | 74.7  |
| 7d      | climatology | 20.09      | 14.60 | 88.8%    | 58.50 | 75.3  |
| 30d     | persistence | 24.88      | 19.63 | 90.3%    | 83.17 | 72.7  |
| 30d     | climatology | 20.93      | 15.23 | 94.6%    | 85.25 | 69.9  |



RMSE by horizon and center (\* = best center; coverage % annotated above each bar). Persistence wins at 24h; climatology overtakes it at 30d on both tracks. Agriculture degrades monotonically with horizon; the air 7d < 1d is a synthetic-generator property (daily random wind dominates the 1-day step but the 7-day traffic weekday cycle aligns), not a real skill gain.

**Phase-1 update (multi-year normals, agriculture,  $\pm 3$  °F on the 60/40 protocol).** Backfilled 3y of NASA POWER daily T2M (data/real\_history\_3y.csv, 17,488 rows) and added a third center — **anomaly persistence**, center =  $\text{normal}(t+h) + \rho(h) \cdot (x_t - \text{normal}(t))$ , with per-site Fourier **ClimateNormals** and  $\rho$  fit per horizon (fitted 1d  $\rho=0.75$ ; 7d/30d  $\rho \rightarrow 0$ ). Best numbers: **1d 88.8%** ( $\pm 1.667$  °C, RMSE 1.131 °C), **7d 71.2%** (RMSE 1.796 °C, skill +0.337 vs climatology), **30d 72.2%** (RMSE 1.802 °C, skill +0.301). Phase-1 exit gate: **30d  $\geq 62\%$  MET; 7d  $\geq 74\%$  missed by 2.8 pts** at the statistical RMSE ceiling ( $\sim 1.80$  °C vs  $\sim 1.50$  needed) — the residual 7-day signal requires NWP input (Phase 2); skill > 0 at every horizon, so Phase 2's MOS has a valid base. Lever sweeps that did not beat per-site harmonics ( $n_{\text{harmonics}}=3$ ): pooled harmonics (7d 63.3%), shared pooled+offsets (7d 64.1%), regional anomalies ( $\rho \rightarrow 0$ ), fixed- $\rho$  0.1 (7d +0.2 pts but 1d 88.8  $\rightarrow$  74.3%),  $n_{\text{harmonics}} \geq 4$  (7d 67.7%).

**Phase-2 update (NWP + MOS, agriculture, ERA5 NWP proxy on the same 60/40 protocol).** MOS center =  $a(h,s) + b(h,s) \cdot \text{NWP}(t+h) + c(h,s) \cdot \text{anomaly}(t)$ , per-site per-horizon OLS with rolling-fit adaptation and

split-conformal 90% intervals; a rolling-skill blend mixes {anomaly, MOS, normal}. NWP here is the Open-Meteo **ERA5 archive** (reanalysis  $\approx$  truth), so these numbers prove the bias-correction machinery, not real forecast skill. Best blend accuracy ( $\pm 3^\circ\text{F}$ ): **1d 98.3%** (RMSE 0.696), **7d 97.3%** (RMSE 0.759 served ensemble), **30d 96.0%** (RMSE 0.793), beating the Phase-1 anomaly center at every horizon; MOS 90% intervals cover 0.86–0.88 (nominal 0.90). The Phase-2 exit gate (7d  $\geq 80\%$ , 30d  $\geq 65\%$ , CRPS(blend) < CRPS(anomaly)) still requires live forecast snapshots accumulated daily (backend open-meteo-source.js is wired for it) or a GFS hindcast archive — blocked on a native GRIB2 decoder on this box.

**Phase-2.5 update (CRPS gate + LightGBM ensemble container).** Analytic Normal CRPS gate CRPS(blend) < CRPS(anomaly) passes at every horizon (1d 0.391 vs 0.606 ... 30d 0.444 vs 0.974). A LightGBM quantile-ensemble container (app/ml/quantile\_ensemble.py, CQR intervals) with the plan's parsimony gate beats the hand blend at  $h \geq 3$  on the proxy (3d 0.746 vs 0.757 ... 30d 0.782 vs 0.793 RMSE) and loses marginally at 1d (0.703 vs 0.696) — so the gate ships the ensemble at  $h \geq 3$  and the blend at 1d. Daily live-forecast snapshots (snapshot\_nwp.py) feed eval\_nwp.py --live once ~3+ weeks accumulate.

**Phase-3 update (conditional conformal calibration, agriculture, 60/40,  $\pm 3^\circ\text{F}$ ).** Per-context conformal widths replace the single per-month qhat: bins of (month  $\times$  seasonal-volatility from the normal's  $\pm 15$ -day std  $\times$  |anomaly| tercile within that month), per-bin finite-sample qhats, month-max fallback, horizon-aware alpha tuned to the top of the coverage band, plus a holdout-calibrated width inflation capped so width targets are never exceeded (eval\_horizons.py --conformal conditional). Gate passes at every horizon — coverage **24h 91.9% / 7d 93.8% / 30d 89.5%** (target band 88–93%) with mean interval width **24h 4.29  $^\circ\text{C}$**  (target  $\leq 4.5$ , was 6.3) and **7d 8.00  $^\circ\text{C}$**  (target  $\leq 8.0$ , was 11.7). The residual-scale analysis shows the 2026 eval window holds genuinely more volatile months (30d Feb |residual| p90 6.17 vs 2.77 calibrated), so the inflation is the honest price within the width budget.

**Phase-4 update (served forecast engine + endpoints).** The Phase-3 stack is packaged for serving: **build\_forecast\_artifact.py** compiles the center hierarchy, per-site ClimateNormals, MOS coefficients, blend weights and the conditional-conformal tables into **models/forecast\_artifact.joblib**, and a new **ForecastEngine** (app/ml/forecast.py) resolves any request site to its nearest grid point and serves it. Lead hierarchy: **h==1 NWP+MOS blend** (anomaly center when no NWP input is supplied), **h==7 P2 quantile ensemble** when NWP is supplied (else NWP+MOS blend, else anomaly center), **h==30 anomaly-persistence**; every interval applies the calibrated per-bin qhat  $\times$  holdout inflation. Exposed as **POST /v1/forecast** (horizon, target date, optional obs/NWP temperature) with model\_version + served\_ts, plus **GET /v1/benchmark** publishing the tolerance-accuracy tables. The backend mirrors both through fail-open clients (ai-benchmark-client.js, /api/benchmark). **Real air labels:** fetch\_openAQ\_labels.py pulls real OpenAQ daily PM2.5 for the 16 sites (OPENAQ\_API\_KEY-gated, synthetic fallback keeps the offline pipeline green) and a CAMS source adapter (cams-source.js) adds a second composition forecast to the air feature group.

**Competitor-style tolerance accuracy** (best center per horizon):


| Horizon | Best center | $\pm 1^\circ\text{C}$ | $\pm 1.667^\circ\text{C}$ | $\pm 2^\circ\text{C}$ | $\pm 3^\circ\text{C}$ |
|---------|-------------|-----------------------|---------------------------|-----------------------|-----------------------|
| 1d      | persistence | 87%                   | 96%                       | 97%                   | 100%                  |
| 7d      | persistence | 44%                   | 68%                       | 77%                   | 90%                   |
| 30d     | anomaly     | 40%                   | 61%                       | 71%                   | 90%                   |

| Horizon | Best center | $\pm 10\ \mu\text{g}/\text{m}^3$ | $\pm 25\ \mu\text{g}/\text{m}^3$ | $\pm 50\ \mu\text{g}/\text{m}^3$ |
|---------|-------------|----------------------------------|----------------------------------|----------------------------------|
| 1d      | anomaly     | 51%                              | 87%                              | 98%                              |
| 7d      | anomaly     | 55%                              | 91%                              | 99%                              |
| 30d     | anomaly     | 45%                              | 83%                              | 97%                              |

Worked example (agriculture, first grid site): **24h** — trained on the observation for **2026-06-21**, predicting **2026-06-22** → center 26.91 °C, actual 26.48 °C, error -0.43 °C. **30d** (best 30-day center, **anomaly**) — fit on observations through **2026-06-03**, predicting **2026-07-03** → center 28.5 °C, actual 30.1 °C, error +1.60 °C. Because the climatology already captures seasonal warming, its 30-day error is smaller than persistence's; the calibrated interval still widens with horizon.

## 7.2 · Competitor Benchmark (tolerance accuracy)


Weather/AQ providers publish **accuracy within tolerance** (typically within  $\pm 3$  °F) rather than calibration scores. This benchmark measures PERN on the same yardstick ( $\pm 3$  °F =  $\pm 1.667$  °C) over the exact 60/40 eval protocol and overlays published figures (NWS/ForecastWatch industry bands, ForecastWatch provider averages, OpenWeather, Weatherbit, Ambee). The Phase-4 headline row is the **served ForecastEngine** — the actual shipped product (h=1 NWP+MOS blend, h=7 P2 quantile ensemble, h=30 anomaly) — scored out-of-sample on the multi-year record.



PERN tolerance-accuracy curves (agriculture) vs published provider numbers at the  $\pm 3$  °F yardstick (dashed vertical). Solid = served engine, dashed = best bare center. Gold squares = published values (ranges as error bars). The curve shows exactly what a single provider % hides: the easy-day mass and the long tail.

### Served engine vs published providers, within $\pm 3$ °F (out-of-sample, multi-year):

| Horizon | Provider (within $\pm 3$ °F) | Published % | PERN served | Verdict     |
|---------|------------------------------|-------------|-------------|-------------|
| 24 h    | NWS / ForecastWatch industry | 90–95%      | 98.3%       | above band  |
| 24 h    | OpenWeather                  | 89%         | 98.3%       | above band  |
| 24 h    | Weatherbit                   | 91%         | 98.3%       | above band  |
| 24 h    | Ambee                        | 90–100%     | 98.3%       | within band |
| 7 d     | NWS / ForecastWatch industry | 70–80%      | 97.3%       | above band  |
| 7 d     | OpenWeather                  | 82%         | 97.3%       | above band  |
| 7 d     | Weatherbit                   | 84%         | 97.3%       | above band  |
| 30 d    | NWS / ForecastWatch industry | 30–60%      | 72.1%       | above band  |






Horizontal bars = published provider bands; green line = PERN served. 24 h 98.3% (NWP+MOS blend; anomaly fallback 88.8% with NWP off) clears every published band on the proxy; 7 d 97.3% clears every provider only on the ERA5 proxy (near-truth, P2 ensemble) — with NWP forced off the same engine serves the anomaly fallback at 71.1%, already inside the NWS 70–80 band; 30 d 72.1% (anomaly, no NWP) clears the NWS 10 d 30–60% band.

### Bare centers (transparency — where the served numbers come from):


| Horizon   | PERN (best center) | PERN within $\pm 3^\circ\text{F}$ | Industry / provider (within $\pm 3^\circ\text{F}$ )                     |
|-----------|--------------------|-----------------------------------|-------------------------------------------------------------------------|
| 24 h      | anomaly            | 88.8%                             | NWS 90–95% · OpenWeather 89% · Weatherbit 91% · Ambee >90%              |
| 7 d       | anomaly            | 71.2%                             | NWS 70–80% · OpenWeather 82% · Weatherbit 84%                           |
| 30 d      | anomaly            | 72.2%                             | NWS ~30–60% (10 d)                                                      |
| 1–5 d avg | —                  | —                                 | ForecastWatch providers: Microsoft 79.5% · TWC/IBM 79.2% · Foreca 77.7% |



Air track tolerance-accuracy curves (PM2.5,  $\mu\text{g}/\text{m}^3$ ). No AQ provider publishes a consistent within-tolerance number, so no competitor markers exist; the curves document PERN's share-within-tolerance at every lead time for honest reporting.

**Honest reading.** PERN scores *daily-mean* temperature while providers score instantaneous/daily-high forecasts — daily means are smoother, so this comparison flatters PERN somewhat. The served 24 h row is the NWP+MOS blend at 98.3% within  $\pm 3^\circ\text{F}$  (94.4% for MOS alone) when an NWP input exists, above every published 24 h band; with NWP forced off the same engine serves the anomaly center at 88.8% (RMSE 1.131, up from the old persistence rung's 87.8%). The 98.3% is proxy-fed (ERA5  $\approx$  truth), so it is not claimable until live-NWP snapshots pass the gate. The honest 7 d row today — the same engine with NWP forced off — is the anomaly fallback at 71.1% within  $\pm 3^\circ\text{F}$ , already inside the NWS 70–80 band; the served 7 d is the P2 quantile ensemble at 97.3% (RMSE 0.759, honest 86.8% coverage, width 5.72) on the proxy, not claimable until ~3+ weeks of live-NWP snapshots pass the gate. Bare anomaly rows use the artifact's anomaly config (harmonic, per-site, 3 harmonics), so the bare 30 d anomaly (72%) equals the served 30 d center (72.1%) exactly. The cold-start 24 h  $\approx 95\%$  figure is a most-recent-month measurement, not the 88.8% multi-year no-NWP headline.

### Cold-start sensitivity — the "user entered 7 days / 2 weeks / a month" case:



Per site, only the last W days of history calibrate, then the same final 30 real days are scored at each lead (best center by RMSE). 24 h is 95% within  $\pm 3^\circ\text{F}$  even with just 7 days of history (persistence needs only today's observation); 7 d sits at the NWS 70%



floor from a week of data (69%) and jumps to 88% once a full seasonal year fits the harmonic anomaly normal; 30 d also needs that full year (62% vs the NWS 30–60 band). The cold-start benchmark uses the artifact's anomaly config (harmonic, per-site). Caveat: the eval window is the single most recent month — it measures the cold-start question, not the full-record 60/40 number (24 h 95% here vs the 88.8% no-NWP served headline). `models/benchmark_history.csv` + `models/charts/benchmark_history_sensitivity.png`.

## 8 · Drift Detection & Safe Retraining

|                            |                                                                                                                                                                                                                                                                                                                                                                             |
|----------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>Feature drift (PSI)</b> | <code>population_stability_index</code> splits the reference distribution into 10 quantile bins and computes $\sum (p_o - p_e) \cdot \ln(p_o/p_e)$ . Moderate > 0.10, severe > 0.25; flagged features appear in the API response.                                                                                                                                           |
| <b>Residual CUSUM</b>      | Two-sided: $s_{\square}(t) = \max(0, s_{\square}(t-1) + r - k\sigma)$ , $s_{\square}$ mirrors, reset on crossing $h\sigma$ . $k = 0.5$ , $h = 5$ in units of reference std — catches a persistent rise in forecast error (stale calibration).                                                                                                                               |
| <b>Serve-time hint</b>     | Every <code>/v1/confidence</code> response embeds drift: <code>max_abs_z</code> and flagged features vs the artifact's <code>reference_stats</code> (training mean/std); $ z  > 4 \rightarrow$ off-distribution.                                                                                                                                                            |
| <b>Retraining gate</b>     | <code>retrain.py</code> rebuilds the dataset, trains a candidate, and replaces <code>models/artifact.joblib</code> only if the candidate temporal backtest is at least as calibrated and accurate as the incumbent. Every decision is logged to <code>models/promotions.jsonl</code> ; <code>--force</code> bypasses the gate. A bad candidate can never reach the website. |

## 9 · Artifact Schema & API Spec

`models/artifact.joblib` (joblib dump of a dict):

```
features: [latitude, longitude, temperature, temperature_max, temperature_min, precipitation,
humidity, wind_speed, month, day_of_year, day_of_week]
model: PersistenceResidual(temp_idx=2)
model_type: persistence · alpha: 0.10 · qhat: 1.261
target_std: 4.009 · target_transform: level
reference_stats: {<feature>: {mean, std}} · trained_ts: ISO-8601
```

POST `/v1/confidence` — request & response:

```
POST /v1/confidence
Request: { latitude: 31.0, longitude: 31.3, feature_group: "agriculture",
features: { temperature: 30.2, temperature_max: 34.1, temperature_min: 26.4,
precipitation: 0.0, humidity: 62, wind_speed: 4.1,
month: 7, day_of_year: 202, day_of_week: 3 },
ts: "2026-07-21T00:00:00Z" }
Response: { score: 87.4, center: 30.2, lower: 28.939, upper: 31.461,
coverage: 0.9, interval_width: 2.5219,
method: "persistence-quantile+cqr", model_version: <trained_ts>,
feature_group: "agriculture", served_ts: <utc>,
drift: { max_abs_z, flagged: [...], hint: "in-distribution" } }
```

## 10 · Code / Module Map

| Module                           | Responsibility                                                                                                                                   |
|----------------------------------|--------------------------------------------------------------------------------------------------------------------------------------------------|
| <code>app/ml/models.py</code>    | LightGBMQuantile, SklearnQuantile, PersistenceResidual, ClimateNormals, DEFAULT_ALPHAS                                                           |
| <code>app/ml/conformal.py</code> | <code>cqr_quantile</code> , <code>cqr_intervals</code> , <code>quantile_residual_intervals</code> , <code>conditional_conformal_intervals</code> |
| <code>app/ml/cv.py</code>        | <code>temporal_block_split</code> , <code>leave_locations_out_split</code> , <code>calibration_split</code>                                      |
| <code>app/ml/metrics.py</code>   | RMSE, MAE, MAPE, pinball/CRPS, coverage, interval width                                                                                          |
| <code>app/ml/backtest.py</code>  | <code>feature_columns</code> , <code>feature_matrix</code> , <code>run_backtest</code> , <code>confidence_score</code>                           |
| <code>app/ml/drift.py</code>     | PSI, <code>feature_psi</code> , <code>residual_cusum</code> , <code>drift_summary</code>                                                         |

| Module                                                                   | Responsibility                                                                                                                                                                                                                                                                                                                             |
|--------------------------------------------------------------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| app/ml/labels.py · synthetic.py                                          | load_training_data (real CSV → synthetic fallback)                                                                                                                                                                                                                                                                                         |
| app/main.py                                                              | FastAPI app: /health, /v1/confidence, /v1/forecast (\$4), /v1/benchmark (\$4)                                                                                                                                                                                                                                                              |
| train.py · retrain.py · backtest_cli.py                                  | train + promote gate · scheduled retrain · CLI backtest                                                                                                                                                                                                                                                                                    |
| make_real_dataset.py · make_real_air_dataset.py · fetch_operaq_labels.py | real POWER labels · synthetic air labels · real OpenAQ daily PM2.5 labels (key-gated, synthetic fallback)                                                                                                                                                                                                                                  |
| eval_horizons.py                                                         | multi-horizon (24h/7d/30d) walk-forward skill eval → models/horizon_eval.json; --conformal conditional = Phase-3 per-context widths (month × vol × [anomaly] bins, tuned alpha, holdout inflation) → models/horizon_eval_conditional.json                                                                                                  |
| app/ml/conformal.py                                                      | grouped + conditional split-conformal intervals (per-bin qhats, month-max fallback)                                                                                                                                                                                                                                                        |
| eval_nwp.py · app/ml/mos.py                                              | Phase-2 NWP+MOS harness: per-site per-horizon OLS bias correction, rolling-fit + split-conformal intervals, rolling-skill blend, CRPS gate, --ensemble parsimony gate → models/phase2_eval.json                                                                                                                                            |
| snapshot_nwp.py · app/ml/quantile_ensemble.py                            | daily Open-Meteo 16-day forecast snapshots → data/nwp_live/ · LightGBM quantile-ensemble container (CQR) for the Phase-2.5 parsimony gate                                                                                                                                                                                                  |
| build_forecast_artifact.py · app/ml/forecast.py                          | Phase-4 serving artifact (center hierarchy + MOS + conditional tables) → models/forecast_artifact.joblib · ForecastEngine: site resolution, lead hierarchy, qhat lookup, /v1/forecast                                                                                                                                                      |
| benchmark_competitors.py                                                 | tolerance-accuracy benchmark vs published providers + month/site strata + served row + cold-start sensitivity (--history-windows) + baseline lock → models/benchmark*.json/csv + charts                                                                                                                                                    |
| check_benchmark.py · run_benchmark_gate.py                               | living-contract gate: fail on regression vs models/benchmark_baseline.json (CI/nightly; retrain.py --benchmark-gate)                                                                                                                                                                                                                       |
| tests/                                                                   | test_backtest (9) · test_api (4) · test_drift (4) · test_conformal_models (7) · test_benchmark_gate (9) · test_mos (6) · test_conformal (6) · test_forecast (7) · test_operaq_labels (6) · test_benchmark_served (4) · test_benchmark_history (4) · test_live_nwp (2) — pytest green (70 passed)                                           |
| pern-backend (Node)                                                      | ai-confidence-client.js + ai-benchmark-client.js (fail-open clients) → trust-engine.js computeConfidenceWithAI(); global-ingestion.js uses the AI path; open-meteo-source.js adds the NWP feature group (live 16-day forecast + ERA5 archive for MOS training); cams-source.js adds the CAMS composition forecast to the air feature group |

## 11 · Technology Stack

| Language / runtime  | Python 3.14.0 (model + API) · Node.js/Express (backend client)                                                                                                                                |
|---------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Web API             | FastAPI 0.141.1 + Uvicorn · Pydantic request validation                                                                                                                                       |
| ML libraries        | LightGBM 4.7.0 (quantile) · scikit-learn 1.9.0 (fallback) · numpy 2.5.2 · pandas 3.0.5 · joblib (artifact serialization)                                                                      |
| Statistical methods | split-conformal CQR · conditional conformal (per-context widths, month × vol × [anomaly] bins) · quantile regression · MOS (rolling OLS bias correction + blend) · PSI · CUSUM                |
| Data sources        | NASA POWER (history) · Open-Meteo (16-day NWP + ERA5 archive for MOS) · OpenAQ (real air labels, key-gated) · CAMS (composition forecast, key-gated, sim fallback)                            |
| Deployment          | Docker — pern-ai service on :8000 wired into docker-compose; artifact hot-swapped by retrain.py without a rebuild; /v1/forecast + /v1/benchmark served from build_forecast_artifact.py output |
| Persistence         | PostgreSQL (feature_vectors, global_api_keys, device keys)                                                                                                                                    |
| Quality gates       | pytest 70 passed · backend vitest 248 passed across 15 files · benchmark gate (check_benchmark.py) fails CI/retrain on regression vs the locked baseline                                      |

## 12 · Honest Limitations

- At a 1-day horizon the served point forecast is the anomaly center (88.8% within  $\pm 3$  °F, up from the old persistence rung's 87.8%); the engine's primary value is calibrated uncertainty, complemented by the NWP+MOS 24 h blend (98.3% on the proxy) once live NWP lands.
- NASA POWER daily values for the most recent 1–2 days are preliminary and can be revised.
- Air quality labels come from OpenAQ once OPENAQ\_API\_KEY is set; without a key the offline pipeline falls back to physically-plausible synthetic PM2.5, so air-track numbers carry that caveat until keyed. CAMS endpoints likewise fall back to simulation without CAMS\_API\_KEY.
- The MOS blend weights were estimated on an ERA5-archive proxy; they are re-estimated from data/nwp\_live/snapshots once ~3 weeks accumulate.
- Spatial leave-locations-out coverage (77%) is lower than temporal — brand-new sites deserve less trust, and the confidence score reflects it.
- The promotion gate compares candidates on the same data family; it prevents regression but cannot invent signal that the features do not contain.